


Mode A (trained to counter Pole A)

Mode B (trained to counter Pole B)



Each panel is a different way the deployment environment can be imperfect (applied only after training, never during it). Under each condition, the SAME frozen policy is tested against two opponent styles, Pole A and Pole B. Good behavior is the blue bar (Mode A) winning more against Pole A and the orange bar (Mode B) winning more against Pole B, in every panel. 128 matched test scenarios per bar. Source: ROBUSTNESS_2V2_RUNG1_RESULT.json.