


Mode A (trained to counter Pole A)

Mode B (trained to counter Pole B)



Each panel is a different way the deployment environment can be imperfect (applied only after training, never during it). Under each condition, the SAME frozen policy is tested against two opponent styles, Pole A and Pole B. Desired specialization is indicated when Mode A outperforms Mode B against Pole A, while Mode B outperforms Mode A against Pole B. Bars show the mean win rate over 128 matched test scenarios; error bars are 95% bootstrap confidence intervals (paired percentile bootstrap, 20,000 resamples). Source: ROBUSTNESS_2V2_RUNG1_RESULT.json.